



CS 655 - Advanced Computer Graphics

Ray Tracing Part I - The Basics

Ray Tracing


What is ray tracing?


- Follow (trace) the path of a ray of light and model how it interacts with the scene
- When a ray intersects an object, send off secondary rays (reflection, shadow, transmission) and determine how they interact with the scene
- Basic algorithm allows for:
 - Hidden surface removal
 - Multiple light sources
 - Hard shadows
 - Reflections
 - Transparent refractions
- Extensions can achieve:
 - Soft shadows
 - Blurred reflections (glossiness)
 - Translucent refractions
 - Motion blur
 - Depth of field (finite apertures)
 - and more

Ray Tracing

- Produces Highly realistic scenes
- Strengths:
 - Specular reflections
 - Transparency
- Weaknesses:
 - Color bleeding (diffuse reflections)
 - Time consuming
- References:
 - “An Improved Illumination Model for Shaded Display,” Turner Whitted, CACM, June 1980.
 - “Distributed Ray Tracing,” Cook, Porter, and Carpenter, Computer Graphics, July 1984, pp. 137-145.

Ray traced images










"Pebbles" by [Jonathan Hunt](#) (2008) 4.5 days to render on an Athlon 5600+




Asbjørn Heid implemented a realistic metal material to render this very realistic buddha model. (from PBRT Gallery)



22.5 fps on a PS3 using 7 cells in 6/2007 by [Eric Rollins](#)


Ray Tracing

- “Backward” ray tracing:
 - Traces the ray *forward* (in time) from the light source through potentially many scene interactions
 - Physically based
 - Global illumination model:
 - Color bleeding
 - Caustics
 - Etc.
 - Problem: most rays will never even get close to the eye
 - Very inefficient since it computes many rays that are never seen



Ray Tracing

- “Forward” ray tracing:
 - Traces the ray *backward* (in time) from the eye, through a point on the screen
 - Not physically based
 - Doesn’t properly model:
 - Color bleeding
 - Caustics
 - Other changes in light intensity and color due to refractions and non-specular reflections
 - More efficient: computes only visible rays (since we start at eye)
 - Generally, “ray tracing” refers to *forward* ray tracing



Ray Tracing

- Ray tracing is an image-precision algorithm: Visibility determined on a per-pixel basis
 - Trace one (or more) rays per pixel
 - Compute closest object (triangle, sphere, etc.) for each ray
- Produces realistic results
- Computationally expensive



1024×1024, 16 rays/pixel
~ 10 hours on a 99 MHz HP workstation

Ray Tracing

- Ray tracing is an image-precision algorithm: **Visibility** determined on a per-pixel basis
 - Trace one (or more) rays per pixel
 - Compute closest object (triangle, sphere, etc.) for each ray
- Produces realistic results
- Computationally expensive



4.5 days on Athlon 5600+

Minimal Ray Tracer

- A basic (minimal) ray tracer is simple to implement:
 - The code can even fit on a 3×5 card (code courtesy of Paul Heckbert with a small change to output as a PPM file):

```
typedef struct{double x,y,z}vec;vec U,black,amb={.02,.02,.02};struct sphere{
vec cen,color;double rad,kd,ks,kt,kl,ir}*s,*best,sph[]={0.,6.,.5,1.,1.,1.,.9,
.05,.2,.85,0.,1.7,-1.,8.,-.5,1.,.5,.2,1.,.7,.3,0.,.05,1.2,1.,8.,-.5,.1,.8,.8,
1.,.3,.7,0.,0.,1.2,3.,-6.,15.,1.,.8,1.,7.,0.,0.,0.,.6,1.5,-3.,-3.,12.,.8,1.,
1.,5.,0.,0.,0.,.5,1.5,};yx;double u,b,tmin,sqrt(),tan();double vdot(A,B)vec A
,B;{return A.x*B.x+A.y*B.y+A.z*B.z;}vec vcomb(a,A,B)double a;vec A,B;{B.x+=a*
A.x;B.y+=a*A.y;B.z+=a*A.z;return B;}vec vunit(A)vec A;{return vcomb(1./sqrt(
vdot(A,A)),A,black);}struct sphere*intersect(P,D)vec P,D;{best=0;tmin=1e30;s=
sph+5;while(s-->sph)b=vdot(D,U=vcomb(-1.,P,s->cen)),u=b*b-vdot(U,U)+s->rad*s
->rad,u=u>0?sqrt(u):1e31,u=b-u>1e-7?b-u:b+u,tmin=u>=1e-7&&u<tmin?best=s,u:
tmin;return best;}vec trace(level,P,D)vec P,D;{double d,eta,e;vec N,color;
struct sphere*s,*l;if(!level--)return black;if(s=intersect(P,D));else return
amb;color=amb;eta=s->ir;d= -vdot(D,N=vunit(vcomb(-1.,P=vcomb(tmin,D,P),s->cen
)));if(d<0)N=vcomb(-1.,N,black),eta=1/eta,d= -d;l=sph+5;while(l-->sph)if((e=l
->kl*vdot(N,U=vunit(vcomb(-1.,P,l->cen))))>0&&intersect(P,U)==l)color=vcomb(e
,l->color,color);U=s->color;color.x*=U.x;color.y*=U.y;color.z*=U.z;e=1-eta*
eta*(1-d*d);return vcomb(s->kt,e>0?trace(level,P,vcomb(eta,D,vcomb(eta*d-sqrt
(e),N,black))):black,vcomb(s->ks,trace(level,P,vcomb(2*d,N,D)),vcomb(s->kd,
color,vcomb(s->kl,U,black))));}main(){puts("P3\n32 32\n255");while(yx<32*32)
U.x=yx%32-32/2,U.z=32/2-yx++/32,U.y=32/2/tan(25/114.5915590261),U=vcomb(255.,
trace(3,black,vunit(U)),black),printf("%.0f %.0f %.0f\n",U);}/*minray!*/
```


Minimal Ray Tracer

- This code implements:
 - Multiple spheres (with different properties)
 - Multiple levels of recursion:
 - Reflections
 - Transparency:
 - Refraction
 - One point light source:
 - Hard shadows
 - Hidden surface removal
 - Phong illumination model
 - It even has a comment



Ray Tracing: Types of Rays

- Primary rays:
 - Sent from the eye, through the image plane, and into the scene
 - May or may not intersect an object in the scene:
 - No intersection \rightarrow set pixel color to background color (P_2)
 - Intersects object \rightarrow send out secondary rays and compute lighting model (P_1)




Ray Tracing: Types of Rays

- Secondary Rays:
 - Sent from the point at which the ray intersects an object
- Multiple types:


Transmission (T): sent in the direction of refraction

Reflection (R): sent in the direction of reflection, and used in the Phong illumination model

Shadow (S): sent toward a light source to determine if point is in shadow or not.




Ray Tracing: Types of Rays



Ray Tracing: Ray Tree

- Each intersection may spawn secondary rays:
 - Rays form a ray tree
 - Nodes → Intersection points
 - Edges → Reflected/transmitted ray
- Rays are recursively spawned until:
 - Ray does not intersect any object
 - Tree reaches a maximum depth
 - Light reaches some minimum value
- Shadow rays are sent from every intersection point (to determine if point is in shadow), but they do not recursively spawn secondary rays

Ray Tracing: Ray Tree Example



Ray tree is evaluated from bottom up:

- Depth-first traversal
- Each node's color is calculated as a function of its children's colors

Basic Ray Tracing Algorithm

- Generate one ray for each pixel
- For each ray:
 - Determine the nearest object intersected by the ray
 - Compute intensity information for the intersection point using the illumination model
 - Calculate and trace reflection ray (if surface is reflective)
 - Calculate and trace transmission ray (if surface is transparent)
 - Calculate and trace shadow ray
 - Combine results of the intensity computation, reflection ray intensity, transmission ray intensity, and shadow ray information
 - If the ray misses all objects, set the pixel color to the background color

Tracing Rays

- Basic (non-recursive) ray tracing algorithm:
 1. Send a ray from the eye through the screen
 2. Determine which object that ray first intersects
 3. Compute pixel color
- Most (approx. 75%) of the time in step 2:
 - Simple method:
 - Compare every ray against every object and remember the closest object hit by each ray
 - Very time consuming:
 - Several optimizations possible

Ray Representation


- A ray can be represented explicitly (in parametric form) as an origin (point) and a direction (vector):

- Origin: $\mathbf{r}_o = \begin{bmatrix} x_o \\ y_o \\ z_o \end{bmatrix}$

- Direction: $\mathbf{r}_d = \begin{bmatrix} x_d \\ y_d \\ z_d \end{bmatrix}$

- The ray consists of all points:


$$\mathbf{r}(t) = \mathbf{r}_o + \mathbf{r}_d t$$



Viewing Ray

- The primary ray (or viewing ray) for a point s on the view plane (i.e., screen) is computed as:

- Origin: $r_o = \text{eye}$
- Direction: $r_d = s - \text{eye}$



- Which coordinate space?
 - Want to define rays in terms world-space coordinates (x, y, z)
 - Eye is already in specified in terms of (x, y, z) position
 - Screen point s is easiest to define in terms of where it is on the window in viewing-space coordinates (u, v, n)

Viewing Ray: Screen Point


- Given:
 - Our scene in world-coordinates
 - A camera (eye) position in world-coordinates (x, y, z)
 - A pixel (i, j) in the viewport
- We need to:
 - Compute the point on the view plane window that corresponds to the (i, j) point in the viewport
 - Transform that point into world-coordinates

Viewport

(i_{\max}, j_{\max})

(i, j)

(i_{\min}, j_{\min})



Viewing Ray (in pictures)

By Convention:

(x,y,z) = world coordinates

(i,j) = viewport coordinates


(u,v,n) = view-reference coordinates

Have s as an (i,j) (inside a dbl for loop)

Need s in (x,y,z) world coordinates


First, take s in (i,j) to (u,v,n) coordinates.

Then take those (u,v,n) coordinates to (x,y,z)



Computing Window Point

- Step 1: Reverse the Window-to-Viewport transformation



Viewport-Window transform

- Window-viewport:

$$i = (u - u_{\min}) \cdot \left(\frac{i_{\max} - i_{\min}}{u_{\max} - u_{\min}} \right) + i_{\min}$$

$$j = (v - v_{\min}) \cdot \left(\frac{j_{\max} - j_{\min}}{v_{\max} - v_{\min}} \right) + j_{\min}$$

- Inverse transform (viewport-window)

$$u = (i - i_{\min}) \cdot \left(\frac{u_{\max} - u_{\min}}{i_{\max} - i_{\min}} \right) + u_{\min}$$

$$v = (j - j_{\min}) \cdot \left(\frac{v_{\max} - v_{\min}}{j_{\max} - j_{\min}} \right) + v_{\min}$$

$$n = 0$$

View-reference to World transform

- Given the screen point in terms of viewing-space coordinates (u, v, n) , transform to world-space (x, y, z) :
 - The viewing transform takes a point from world space to view space:


$$\mathbf{M}_v = \mathbf{RT} = \begin{bmatrix} u_x & u_y & u_z & 0 \\ v_x & v_y & v_z & 0 \\ n_x & n_y & n_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -LookAt_x \\ 0 & 1 & 0 & -LookAt_y \\ 0 & 0 & 1 & -LookAt_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- We want to reverse this:

$$\mathbf{s} = \mathbf{M}_v^{-1} \begin{bmatrix} u_s \\ v_s \\ n_s \\ 1 \end{bmatrix} = \mathbf{T}^{-1} \mathbf{R}^{-1} \begin{bmatrix} u_s \\ v_s \\ n_s \\ 1 \end{bmatrix} = \begin{bmatrix} u_x & v_x & n_x & LookAt_x \\ u_y & v_y & n_y & LookAt_y \\ u_z & v_z & n_z & LookAt_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} u_s \\ v_s \\ n_s \\ 1 \end{bmatrix}$$

or

$$\mathbf{s} = \mathbf{LookAt} + u_s \mathbf{u} + v_s \mathbf{v} + n_s \mathbf{n}$$



M_v = world to view (given x,y,z return u,v,n)


$$\mathbf{M}_v = \mathbf{RT} = \begin{bmatrix} u_x & u_y & u_z & 0 \\ v_x & v_y & v_z & 0 \\ n_x & n_y & n_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -LookAt_x \\ 0 & 1 & 0 & -LookAt_y \\ 0 & 0 & 1 & -LookAt_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Translate lookAt point to origin


“Rotate” u,v,n axes to line up with x,y,z axes.

M_v^{-1} = view to world (given u,v,n return x,y,z)


$$\mathbf{M}_v^{-1} = \mathbf{T}^{-1}\mathbf{R}^{-1} = \begin{bmatrix} u_x & v_x & n_x & LookAt_x \\ u_y & v_y & n_y & LookAt_y \\ u_z & v_z & n_z & LookAt_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



Example



Example



Given:

eye position = 5,4,3, look-at = 10,7,2 and camera-up = 0,1,0

compute (in world coordinates)

n (world) = from - at = $(5,4,3) - (10,7,2) = (-5,-3,1)$

u (world) = up \times n = $(0,1,0) \times (-5,-3,1) = (1,0,5)$

v (world) = $n \times u = (-5,-3,1) \times (1,0,5) = (-15,26,3)$


$$n = 0$$


$$u = (i - i_{\min}) \cdot \left(\frac{u_{\max} - u_{\min}}{i_{\max} - i_{\min}} \right) + u_{\min}$$

$$v = (j - j_{\min}) \cdot \left(\frac{v_{\max} - v_{\min}}{j_{\max} - j_{\min}} \right) + v_{\min}$$

$$n = 0$$

Example: $(i,j) \rightarrow (u,v,n)$

$n(\text{world}) = (-5, -3, 1)$
 $u(\text{world}) = (1, 0, 5)$
 $v(\text{world}) = (-15, 26, 3)$
 $s(\text{view-ref}) = (0.218, -0.152, 0)$



Compute s in world coordinates.

$$s(\text{world}) = M_v^{-1} s(\text{view-ref})$$

$$\begin{bmatrix} 1 & -15 & -5 & 10 \\ 0 & 26 & -3 & 7 \\ 5 & 3 & 1 & 2 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0.218 \\ -0.152 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 10.11 \\ 6.86 \\ 2.19 \\ 1 \end{bmatrix}$$